ANTHROPIC STUDY, OCT 2026

Expert analysts followed a coin flip 86% of the time

Anthropic secretly set which side its AI agent argued for. 50 analysts using Claude Code picked that side 86% of the time. 8 slides - save for your next agent review process

THE SETUP

01. 50 data analysts, one business decision

They used Claude Code to judge if a farming programme raised yields enough to expand it. Four method choices. Each had two defensible answers.

THE COIN FLIP

02. A hidden prompt told the agent which side to argue

One side per choice, set at random. The agent argued honestly and could change its mind if the evidence demanded.

RESULT 1

86%

Analysts matched the agent's hidden side. Chance would be about 50%. The agent framed each choice before the analyst had formed a view, so review turned into agreement.

RESULT 2

Even analysts who wrote their view first switched

AGENT DISAGREED: 60% switched to the agent's side. AGENT AGREED: 6% switched

RESULT 3

They did not notice

77% repeated the agent's arguments as their own. Only 15% credited the agent.

RECAP

What this means for 'a person reviews it'

1. Review after the agent speaks turns into agreement. 2. Make the agent ask for the reviewer's view first. 3. Show the best case for each option, then advise. 4. Flag which judgement calls move the final answer.

Human review is not a control unless the human speaks first.

Save this for your next agent review process.